

446 Section 4 $\leftarrow (3 - \eta(-1))$

TA: Varun Ananth

Plans for today!

1. This
2. Reminders
3. Train/Val/Test Problems
4. Gradient Descent
5. Generalized Least Squares
6. Ridge/LASSO (if time)

Reminders

- HW1 was due yesterday
 - Remember the late day policy!
- HW2 is released
- Midterm in a week...
 - February 7th, Friday


Problems 1.1, 1.2

You are given blocks of code, and something is wrong/not totally right with how they deal with the data.

Identify them and propose solutions!

What do you never ever ever ever ever ever ever do?

Tune your model on your test set!



1.1. Program 1

```
1 # Given dataset of 1000-by-50 feature
2 # matrix X, and 1000-by-1 labels vector
3
4 mu = np.mean(X, axis=0)
5 X = X - mu
6
7 idx = np.random.permutation(1000)
8 TRAIN = idx[0:900]
9 TEST = idx[900::]
10
11 ytrain = y[TRAIN]
12 Xtrain = X[TRAIN, :]
13
14 # solve for argmin_w ||Xtrain*w - ytrain||_2
15 w = np.linalg.solve(np.dot(Xtrain.T, Xtrain), np.dot(Xtrain.T, ytrain))
16
17 b = np.mean(ytrain)
18
19 ytest = y[TEST]
20 Xtest = X[TEST, :]
21
22 train_error = np.dot(np.dot(Xtrain, w)+b - ytrain,
23                      np.dot(Xtrain, w)+b - ytrain ) / len(TRAIN)
24 test_error = np.dot(np.dot(Xtest, w)+b - ytest,
25                      np.dot(Xtest, w)+b - ytest ) / len(TEST)
26
27 print('Train error = ', train_error)
28 print('Test error = ', test_error)
```

mu is calculated from the entire data (train + test), intertwining them!

This is bad!

Calculate a mean just on the train data, and use this to de-mean both the train and test datasets

1.2. Program 2

```
1 # We are given: 1) dataset X with n=1000 samples and 50 features and 2) a vector y of length 1000 with labels.
2 # Consider the following code to train a model, using cross validation to perform hyperparameter tuning.
3
4 def fit(Xin, Yin, _lambda):
5     w = np.linalg.solve(np.dot(Xin.T, Xin) + _lambda * np.eye(Xin.shape[1]), np.dot(Xin.T, Yin))
6     b = np.mean(Yin) - np.dot(w, mu)
7     return w, b
8
9 def predict(w, b, Xin):
10    return np.dot(Xin, w) + b
11
12 idx = np.random.permutation(1000)
13 TRAIN = idx[0:800]
14 VAL = idx[800:900]
15 TEST = idx[900::]
16
17 ytrain = y[TRAIN]
18 Xtrain = X[TRAIN, :]
19 yval = y[VAL]
20 Xval = X[VAL, :]
21
22 # demean data
23 mu = np.mean(Xtrain, axis=0)
24 Xtrain = Xtrain - mu
25 Xval = Xval - mu
26
27 # use validation set to pick the best hyper-parameter to use
28 lambdas = [10 ** -5, 10 ** -4, 10 ** -3, 10 ** -2]
29 err = np.zeros(len(lambdas))
30
31 for idx, _lambda in enumerate(lambdas):
32     w, b = fit(Xtrain, ytrain, _lambda)
33     yval_hat = predict(w, b, Xval)
34     err[idx] = np.mean((yval_hat - yval)**2)
35
36 lambda_best = lambdas[np.argmin(err)]
37
38 Xtot = np.concatenate((Xtrain, Xval), axis=0)
39 ytot = np.concatenate((ytrain, yval), axis=0)
40
41 w, b = fit(Xtot, ytot, lambda_best)
42
43 ytest = y[TEST]
44 Xtest = X[TEST, :]
45
46 # demean data
47 Xtest = Xtest - mu
48
49 ytot_hat = predict(w, b, Xtot, lambda_best)
50 train_error = np.mean((ytot_hat - ytot)**2)
51 ytest_hat = predict(w, b, Xtest, lambda_best)
52 test_error = np.mean((ytest_hat - ytest)**2)
53
54 print('Train error = ', train_error)
55 print('Test error = ', test_error)
```

The final model is trained on BOTH the training and validation sets.

This is... eh...

Your hyperparameters selected on just the train data may not hold for train + val

- Tradeoff between more data and better test error estimate




Gradient Descent

Gradient Descent

Purpose of this exercise:
Understanding how gradient descent relates to approximations, and why it works.

Consider some function $f(w)$, which has some w_* for which $w_* = \arg \min_w f(w)$:



2a

- (a) For some w that is very close to w_0 , give the Taylor series approximation for $f(w)$ starting at $f(w_0)$.

Remember Taylor expansion?

↳ To approximate a function around a point \underline{a}

$$f(x) \approx f(a) + \frac{f'(a)}{1!}(x-a) + \frac{f''(a)}{2!}(x-a)^2 + \frac{f'''(a)}{3!}(x-a)^3 \dots$$


↑
Exact at \underline{a} , close around a

Better and better approximations

2a (answer)


- (a) For some w that is very close to w_0 , give the Taylor series approximation for $f(w)$ starting at $f(w_0)$.

For w very close to w_0 , we see that $f(w) \approx f(w_0) + (w - w_0) \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right)$.



2b

- (b) Now, let us choose some $\eta > 0$ that is *very small*. With this very small η , let's assume that $w_1 = w_0 - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right)$. Using your approximation from part (a), give an expression for $f(w_1)$.


$$f(w) \approx f(w_0) + (w - w_0) \underbrace{\left(\frac{df(w)}{dw} \Big|_{w=w_0} \right)}.$$

Hint: Plug in
here

Fancy way of saying $f'(w_0)$

(Derivative of $f(w)$ at w_0)

2b (answer)

$$w_1 = w_0 - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right) \leftarrow \text{Given}$$

$$\begin{aligned} f(w_1) &\approx f(w_0) + (w_1 - w_0) \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right) \\ &= f(w_0) + \left(w_0 - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right) - w_0 \right) \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right) \\ &= f(w_0) - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right)^2 \end{aligned}$$

2c

- (c) Given your expression for $f(w_1)$ from part (b), explain why, if η is small enough and if the function approximation is a good enough approximation, we are guaranteed to move in the “right” direction closer to the minimum w_* .

Remember:

We want to
minimize this

$$f(w_1) \approx f(w_0) - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right)^2$$

Hint: Why
would this be
good?

2c (answer)

Note that in part (b), the derivative is squared and will always be a nonnegative value. Therefore, $f(w_1) < f(w_0)$.

$$f(w_1) \approx f(w_0) - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_0} \right)^2$$

In English: The loss function after a weight update will always evaluate to be smaller than before the weight update

- If the step size is small enough
- If the approximation is good enough

2d (answer)


Gradient descent is written as:

$$\text{For } k = 0, 1, 2, 3, \dots, w_{k+1} = w_k - \eta \left(\frac{df(w)}{dw} \Big|_{w=w_k} \right).$$

$$\text{Note that as } k \rightarrow \infty, \left(\frac{df(w)}{dw} \Big|_{w=w_k} \right) \rightarrow 0.$$

Convergence guarantees iff convex!

We visualize as:



Generalized Least Squares

Least Squares Proof(s)

Has shown up...

- In lecture (Lecture 2)
- On your homework
(A5 Ridge
Regression proof)
- And now here!

You can look at the
generalized proof in your
own time.

Should look familiar...

$$\hat{\omega}_{\text{general}} = (X^\top X + \lambda D)^{-1} X^\top y$$

$$\hat{\omega}_{\text{general}} = \left(\sum_{i=1}^n x_i x_i^\top + \lambda D \right)^{-1} \left(\sum_{i=1}^n x_i y_i \right)$$

3.2a

$$\hat{\omega}_{\text{general}} = (X^\top X + \lambda D)^{-1} X^\top y$$

- (a) In the simple least squares case ($\lambda = 0$ above), what happens to the resulting $\hat{\omega}$ if we double all the values of y_i ?

3.2a (answer)

$$\hat{\omega}_{\text{general}} = (X^\top X + \lambda D)^{-1} X^\top y$$

- (a) In the simple least squares case ($\lambda = 0$ above), what happens to the resulting $\hat{\omega}$ if we double all the values of y_i ?

Solution:

As can be seen from the formula $\hat{\omega} = (X^\top X)^{-1} X^\top y$, doubling y doubles $\hat{\omega}$ as well. This makes sense intuitively as well because if the observations are scaled up, the model should also be.

3.2b

$$\widehat{\omega}_{\text{general}} = (X^\top X + \lambda D)^{-1} X^\top y$$

- (b) In the simple least squares case ($\lambda = 0$ above), what happens to the resulting $\widehat{\omega}$ if we double the data matrix $X \in \mathbb{R}^{n \times d}$?

3.2b (answer)

$$\hat{\omega}_{\text{general}} = (X^\top X + \lambda D)^{-1} X^\top y$$


- (b) In the simple least squares case ($\lambda = 0$ above), what happens to the resulting $\hat{\omega}$ if we double the data matrix $X \in \mathbb{R}^{n \times d}$?

Solution:


As can be seen from the formula $\hat{\omega} = (X^\top X)^{-1} X^\top y$, doubling X halves ω . This also makes sense intuitively because the error we are trying to minimize is $\|X\omega - y\|_2^2$, and if the X has doubled, while y has remained unchanged, then ω must compensate for it by reducing by a factor of 2.

Ridge vs. LASSO

ridge regularization (L2)



LASSO regularization (L1)



Questions/Chat Time!