

Practice 2: PyTorch Basics and Model Training

Outline

- Neural Net Training Workflow
- Pytorch Data Type: Tensors
- Graph Computation and Neural Net Models
- Example: Iris Dataset Classification
- Assignment: MNIST Classification

Part 1: Neural Net Training Workflow in Pytorch

Neural Net Workflow Steps

- Prepare Data
- Select Hyperparameter
- Define Model
- Identify Tracked Values
- Train and Validate Model
- Visualization and Evaluation

Neural Net Workflow Steps

- Prepare Data
 - Select Hyperparameter
 - Define Model
 - Identify Tracked Values
 - Train and Validate Model
 - Visualization and Evaluation
-
- Data Preparation
 - Define batch size
 - Split train/val/test sets
 - Migrate to Tensors
 - Additional pre-processing
(normalization, encoding, etc.)
 - One-hot encoding?

Data-Preprocessing: One-hot Encoding

- One-hot encoding transforms categorical data into one-hot vectors
 - 1 in the index representing your class
 - 0 in all other indices
 - Total dimensions goes from 1 to number of classes
- Use `torch.nn.functional.one_hot()`

Human-Readable

Pet
Cat
Dog
Turtle
Fish
Cat

Machine-Readable

Cat	Dog	Turtle	Fish
1	0	0	0
0	1	0	0
0	0	1	0
0	0	0	1
1	0	0	0

Neural Net Workflow Steps

- Prepare Data
- Select Hyperparameter
- Define Model
- Identify Tracked Values
- Train and Validate Model
- Visualization and Evaluation
- Hyperparameter Selection
 - Network size and type
 - Learning Rate
 - Regularizers and strength
 - Define loss function and optimizer
 - Other hyperparameters

Neural Net Workflow Steps

- Prepare Data
- Select Hyperparameter
- Define Model
 - Identify Tracked Values
 - Train and Validate Model
 - Visualization and Evaluation
- Model Definition
 - Network Type
 - Network Parameters/Layers
 - Output value(s) and dimensions
 - Forward() Function

Neural Net Workflow Steps

- Prepare Data
 - Select Hyperparameter
 - Define Model
 - Identify Tracked Values
 - Train and Validate Model
 - Visualization and Evaluation
- Values to Track
 - Training Loss
 - Validation Loss
 - Other relevant values
 - Create blank placeholders to populate during training
 - Empty lists, arrays, or tensors

Neural Net Workflow Steps

- Prepare Data
- Select Hyperparameter
- Define Model
- Identify Tracked Values
- Train and Validate Model
- Visualization and Evaluation
- Train Model
 - Calculate loss on training set
 - Backpropagate gradients
 - Update weights
- Validate Model
 - Calculate error on validation or test set
 - Do not update weights
- Save losses in placeholders

Neural Net Workflow Steps

- Prepare Data
- Select Hyperparameter
- Define Model
- Identify Tracked Values
- Train and Validate Model
- Visualization and Evaluation
- Visualize Training Progress
 - Plot training and validation losses over course of training
 - Do curves converge?
 - Does loss go up over time?
- Evaluate model
 - Confusion matrix
 - Generate samples
 - Identify model weaknesses

Part 2: Tensors

Torch Tensors

- Main data structure for PyTorch
- Like numpy arrays, but optimized for machine learning
 - Can be migrated to or stored on GPUs
 - Optimized for automatic differentiation
- Three main attributes:
 - Shape – size of each dimension
 - Datatype – form of each entry (float, int, etc.)
 - Device – cpu or cuda (gpu)

Tensor Initialization

- Can create tensor from existing data
 - `torch.Tensor([[1, 2], [3,4]])`
 - `torch.Tensor(np_array)`
- Can generate tensor with random or fixed values
 - `torch.ones(shape)`
 - `torch.zeros(shape)`
 - `torch.rand(shape)`

Tensor Operations


Most numpy operations are identical for tensors

- Indexing and slicing
 - e.g. `tensor[:2, 3:5]` selects the first 2 rows and 4th and 5th column of tensor
- Elementwise addition, subtraction, multiplication
- Matrix/tensor multiplication:
`tensor1.matmul(tensor2)` OR
`tensor1@tensor2`
 - If tensors have 3 dimension, the first dimension will be treated as the batch (so matmul will be conducted on the other 2 dimensions)

$$\begin{matrix} n \times d_1 \times d_2 \\ @ \\ n \times d_2 \times d_3 \end{matrix} = \begin{matrix} n \times d_1 \times d_3 \end{matrix}$$

Tensor Operations

- Concatenate tensors using `torch.cat()`
 - Input should be list or tuple of tensors
 - Can specify dimension over which concatenation occurs (using `dim = ?`)
 - All other dimensions must be identical



Part 3: Graph Computation and Neural Net Models

Neural Net models in PyTorch

- Base class: nn.Module
- Two primary features of base class:
 - Parameters
 - Forward
- Common PyTorch Layers:
 - Linear
 - Activation Functions (ReLU, tanh, etc.)
 - Dropout
 - RNN
 - Convolution

Neural Network Models

```
1 class ExModel(nn.Module):
2     """
3         Example model: A simple, two-layer NN with ReLU activation.
4         Input dimension is 100, output dimension is 1.
5     """
6     def __init__(self):
7         super(ExModel, self).__init__()
8         self.layer1 = nn.Linear(100, 50)
9         self.act1 = nn.ReLU()
10        self.output_layer = nn.Linear(50, 1)
11
12    def forward(self, x):
13        x = self.layer1(x)
14        x = self.act1(x)
15        output= self.output_layer(x)
16        return output
17
18
19 model = ExModel()
20 print(model)
```

- Initialization


- Define layers, parameters of your network
- The parameters of all layers in the network are included in the parameters of the overall network

- Forward()

- Defines how the network processes input- called using `model(input)`. Use layers/parameters defined above
- Never use `model.forward(input)`

Computational graphs

- PyTorch generates a computational graph every time a parameter or variable with `requires_grad` is operated on
- The graph is used to back-propagate errors and update the parameters




Loss functions and Optimizers

- The loss function defines how you penalize errors
 - nn.MSELoss() for regression
 - nn.NLLLoss() or nn.CrossEntropy() for classification
- Optimizer defines how you update the parameter weights given the gradients given by the loss
 - Stochastic Gradient Descent (optim.SGD)
 - Momentum-based, adaptive approach- Adam (optim.Adam)
- Optimizer needs to be given the learning rate and the model parameters to update

```
1 optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```


Optimization during training

- Each training step, the optimization occurs in 3 steps
 - `optimizer.zero_grad()` – Resets the accumulated gradients
 - `loss.backward()` – Back-propagates the gradients to assign the contribution from each parameter
 - `optimizer.step()` – Updates the parameters based on the gradients, according to the optimization scheme (optimizer)




Saving and Loading Models

- Useful quantities to track during training:
 - Training Loss
 - Validation Loss
 - Model state dictionary (parameters)
 - Optimizer state dictionary
- Loading state dictionaries
 - Model: `model.load_statedict()` loads the saved parameter weights into the model
 - Optimizer: `optimizer.load_statedict()` loads optimizer state, such as learning rate, momentum, etc.



Saving and Loading Models

- Can use `torch.save()` and `torch.load()`
 - Operate similarly to `pkl.dump` and `pkl.load`
- Create checkpoints to save all in one file
- Can save every epoch, or define some condition for saving (best loss, every n epochs, etc.)



```
ckpt = {'train_losses': train_losses, 'model_weights': model.state_dict(),
        'optimizer_state': optimizer.state_dict()}
```

Example: Iris Dataset Classification

Tutorial adapted from: <https://janakiev.com/blog/pytorch-iris/>

Iris Dataset


- Dataset containing characteristics of 3 different flower types
- Found in `sklearn.datasets`
- Preprocess data:
 - Train Test split
 - Standard Scaler (Normalize)



```
1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4
5 iris = load_iris()
6 X = iris['data']
7 y = iris['target']
8 names = iris['target_names']
9 feature_names = iris['feature_names']
10
11 # Scale data to have mean 0 and variance 1
12 # which is importance for convergence of the neural network
13 scaler = StandardScaler()
14 X_scaled = scaler.fit_transform(X)
15
16 # Split the data set into training and testing
17 X_train, X_test, y_train, y_test = train_test_split(
18     X_scaled, y, test_size=0.2, random_state=2)
```

Iris Dataset

```
1 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
2 for target, target_name in enumerate(names):
3     X_plot = X[y == target]
4     ax1.plot(X_plot[:, 0], X_plot[:, 1],
5               linestyle='none',
6               marker='o',
7               label=target_name)
8     ax1.set_xlabel(feature_names[0])
9     ax1.set_ylabel(feature_names[1])
10    ax1.axis('equal')
11    ax1.legend();
12
13 for target, target_name in enumerate(names):
14     X_plot = X[y == target]
15     ax2.plot(X_plot[:, 2], X_plot[:, 3],
16               linestyle='none',
17               marker='o',
18               label=target_name)
19     ax2.set_xlabel(feature_names[2])
20     ax2.set_ylabel(feature_names[3])
21     ax2.axis('equal')
22     ax2.legend();
```



- Data Visualization
 - Plotting features against each other
 - Not linearly separable

Model Definition

- 3-layer fully connected neural net
- 1st and 2nd layer have 50 neurons each
- Final layer has 3 neurons
 - Used for classification
- Activation functions can also be called using torch.nn.functional

```
1 import torch.nn.functional as F
2 class Model(nn.Module):
3     def __init__(self, input_dim):
4         super(Model, self).__init__()
5         self.layer1 = nn.Linear(input_dim, 50)
6         self.layer2 = nn.Linear(50, 50)
7         self.layer3 = nn.Linear(50, 3)
8
9     def forward(self, x):
10        x = F.relu(self.layer1(x))
11        x = F.relu(self.layer2(x))
12        x = F.softmax(self.layer3(x), dim=1)
13        return x
```

Data, hyperparameters, and Saved Values

- Define loss and optimizer
- Define training epochs and data
- Tracking loss and accuracy at each epoch

```
1 model      = Model(X_train.shape[1])
2 optimizer  = torch.optim.Adam(model.parameters(), lr=0.001)
3 loss_fn    = nn.CrossEntropyLoss()
```

```
1 import tqdm
2
3 EPOCHS   = 100
4 X_train = torch.from_numpy(X_train).float()
5 y_train = torch.from_numpy(y_train).long()
6 X_test  = torch.from_numpy(X_test).float()
7 y_test  = torch.from_numpy(y_test).long()
8
9 loss_list    = np.zeros((EPOCHS,))
10 accuracy_list = np.zeros((EPOCHS,))
```

Model Training and validation

- Output of model gives predictions
- Track training loss as loss
- Check accuracy on validation set

```
for epoch in tqdm.trange(EPOCHS):
    y_pred = model(X_train)
    loss = loss_fn(y_pred, y_train)
    loss_list[epoch] = loss.item()

    # Zero gradients
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    with torch.no_grad():
        y_pred = model(X_test)
        correct = (torch.argmax(y_pred, dim=1) == y_test).type(torch.FloatTensor)
        accuracy_list[epoch] = correct.mean()
```

Plot Validation accuracy and training loss

```
1 fig, (ax1, ax2) = plt.subplots(2, figsize=(12, 6), sharex=True)
2
3 ax1.plot(accuracy_list)
4 ax1.set_ylabel("validation accuracy")
5 ax2.plot(loss_list)
6 ax2.set_ylabel("training loss")
7 ax2.set_xlabel("epochs");
```

- Loss should go down
- Accuracy should go up

