

Practice 3: Model Optimization and HyperParameter Tuning

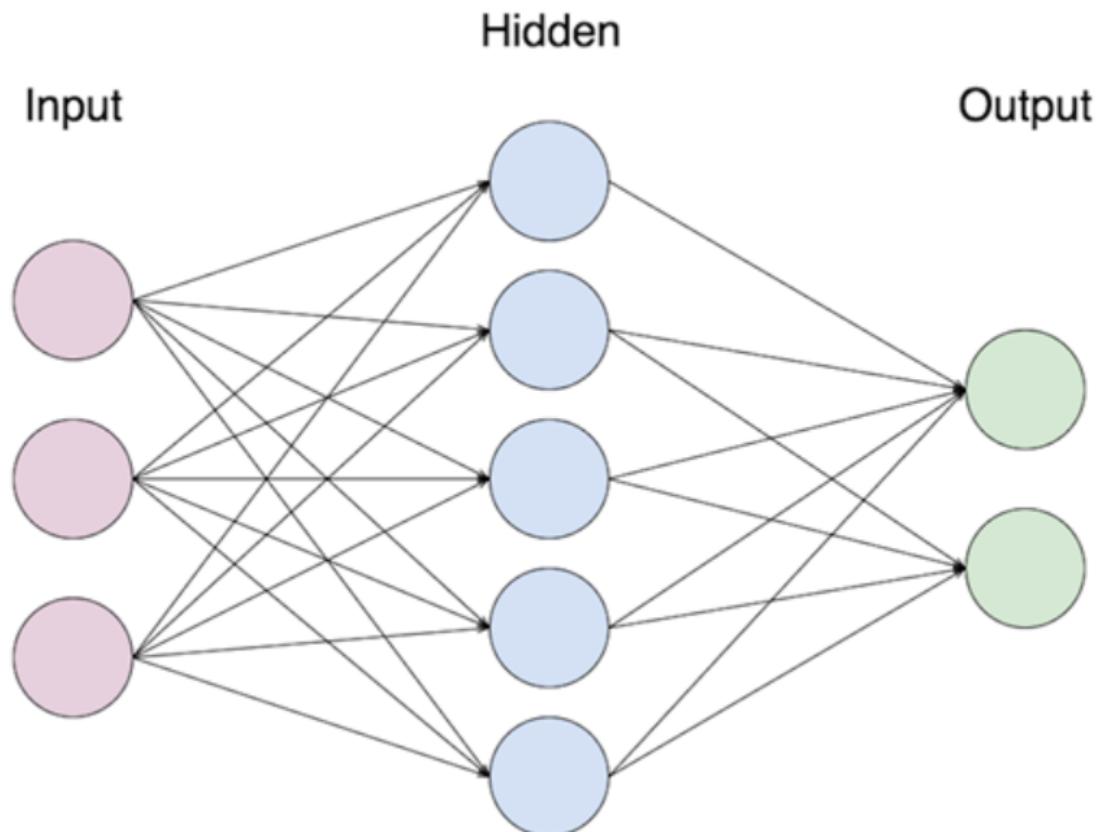
Outline

- Python Operators
 - Activation Functions
 - Normalization
 - Dropout
- Designing Training Procedures

PyTorch Operators and Layers

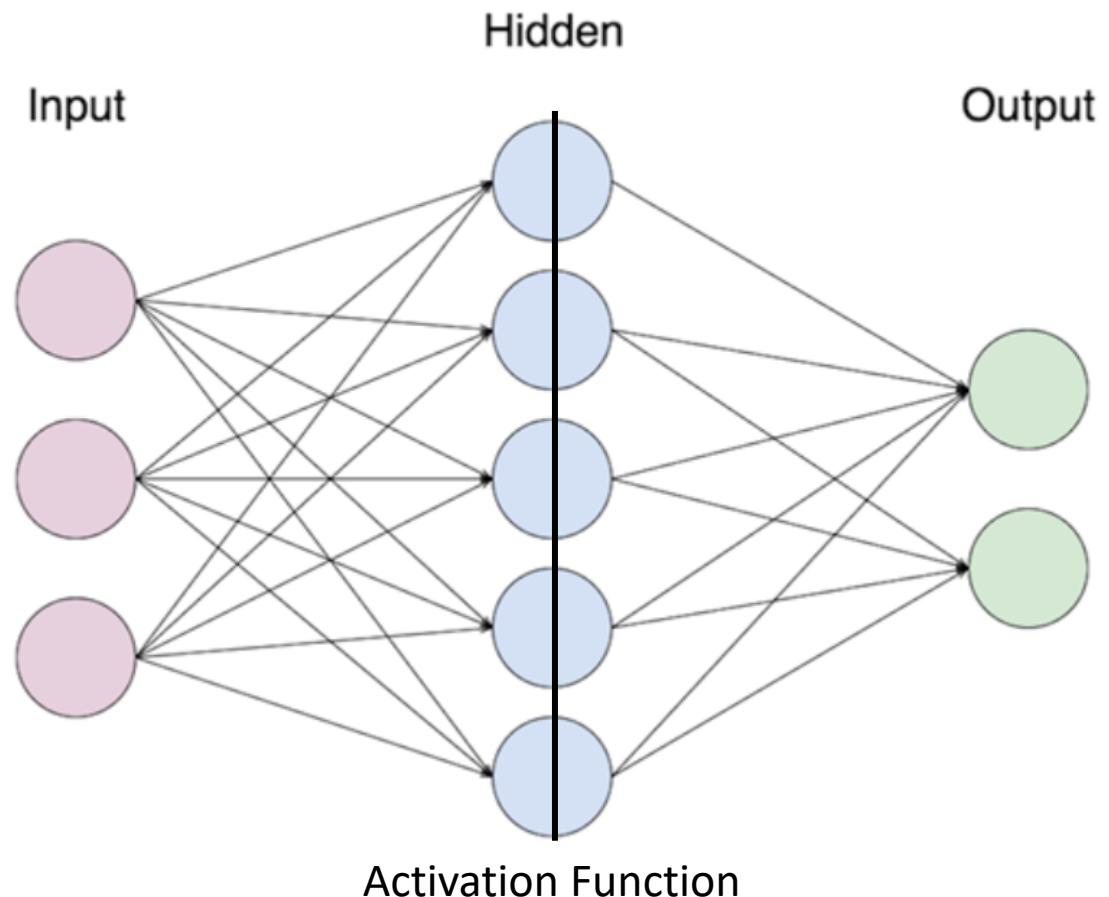
PyTorch Operators/Layers

- Activation Functions
- Normalization
- Dropout



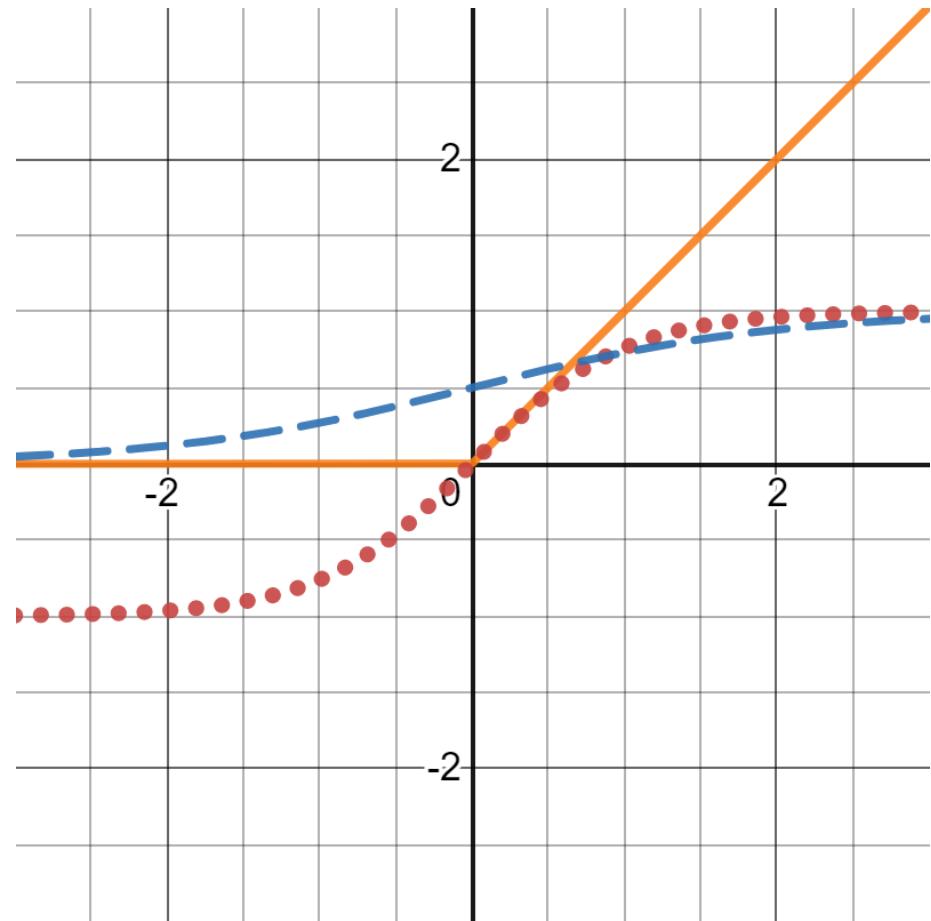
PyTorch Operators/Layers

- Activation Functions
- Normalization
- Dropout
- Initialization



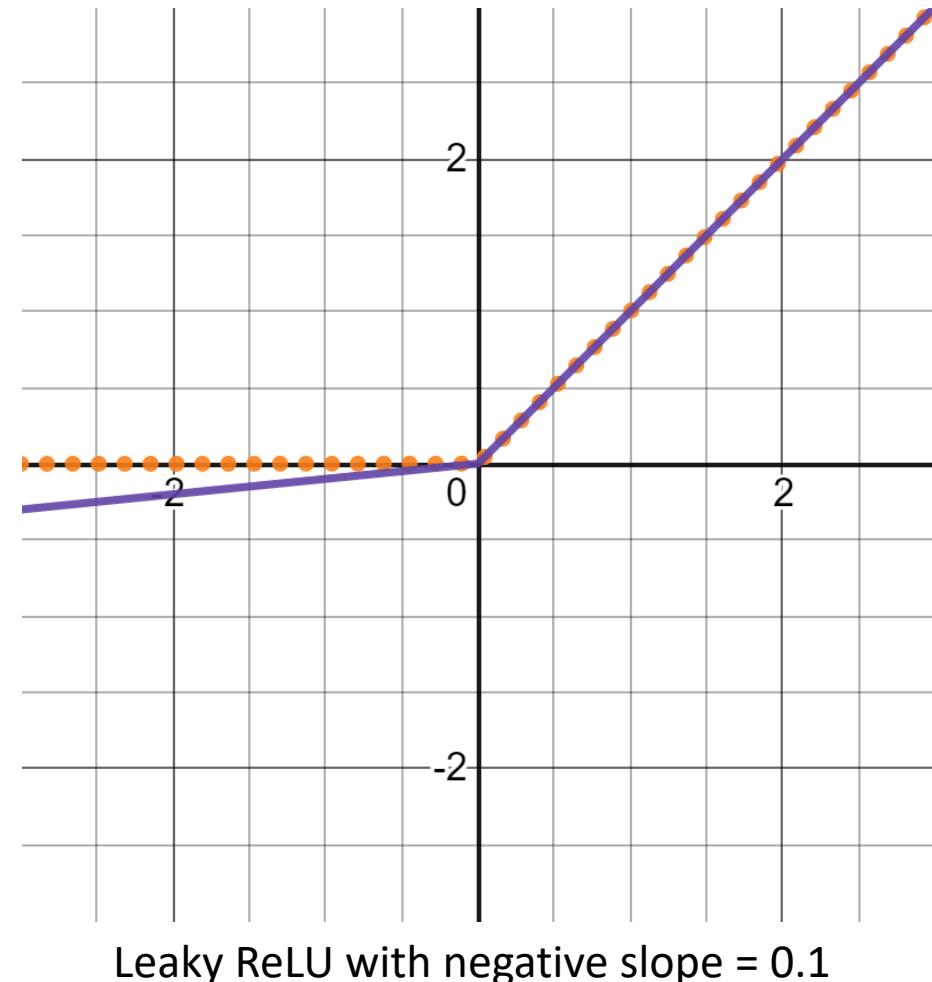
Activation Functions

- Non-linear functions performed by neurons
- ReLU - Rectified Linear Unit (nn.ReLU)
 - $y \geq 0$
- Tanh (nn.tanh)
 - $-1 < y < 1$
 - nn.Tanh
- Sigmoid (nn.Sigmoid)
 - $0 < y < 1$



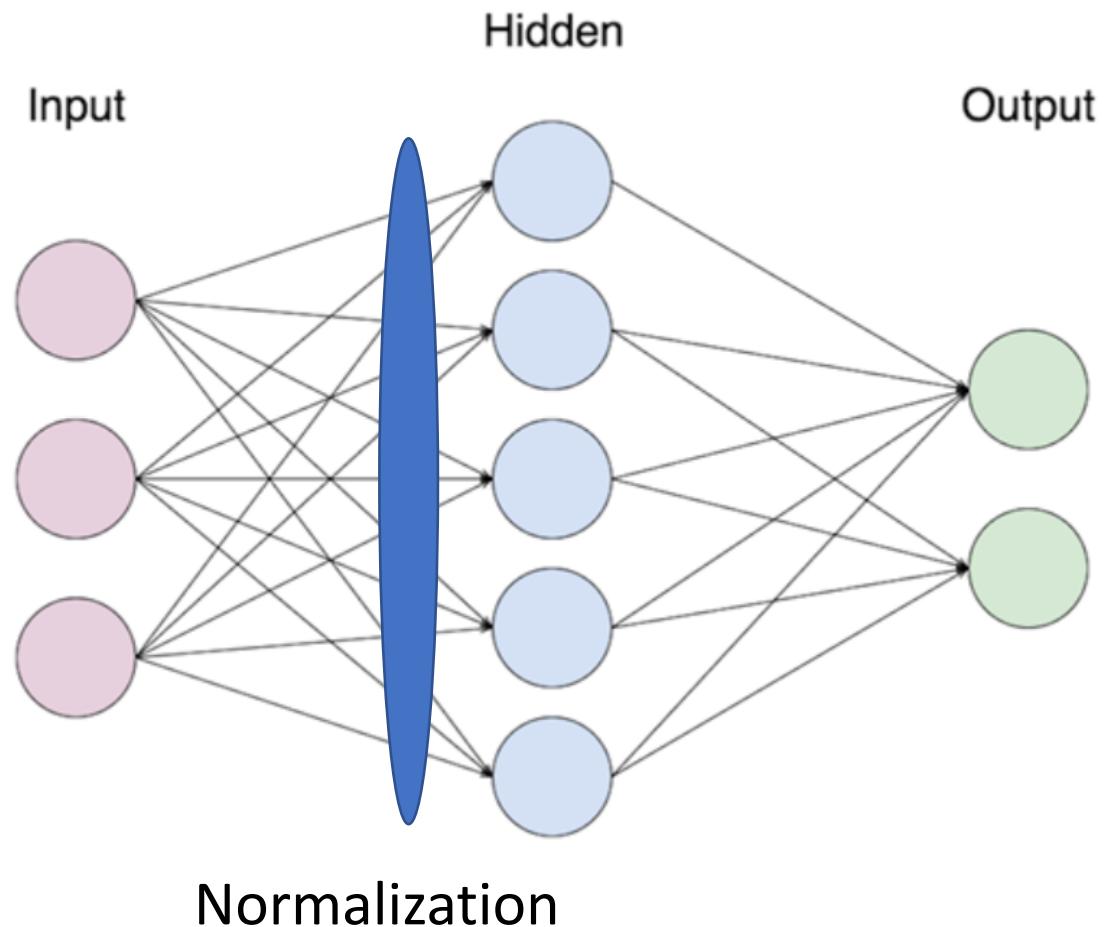
Activation Functions

- Leaky ReLU
 - Similar to ReLU, but has non-zero values for negative x
 - Takes argument *negative_slope*, which determines the slope for $x < 0$.
- For full list of activation functions, see:
<https://pytorch.org/docs/stable/nn.html>



Pytorch Operators/Layers

- Activation Functions
- Normalization
- Dropout

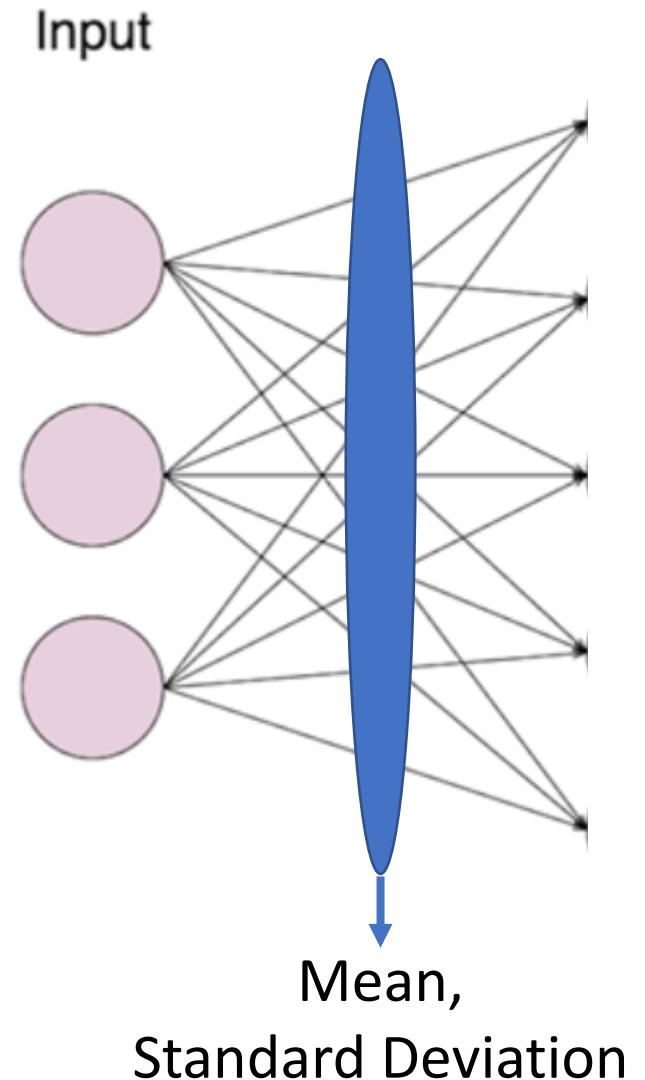


Input Normalization: Batch Normalization

- Normalizes input into each layer for each training mini-batch
- Addresses issue of shifting input distributions over training
- Inputs:
 - num_features: Number of features in the input vector
 - eps: numerical stability parameter

Syntax Example:

```
self.bn2 = torch.nn.BatchNorm1D(input_dim)  
self.layer2 = torch.nn.Linear(input_dim, output_dim)
```



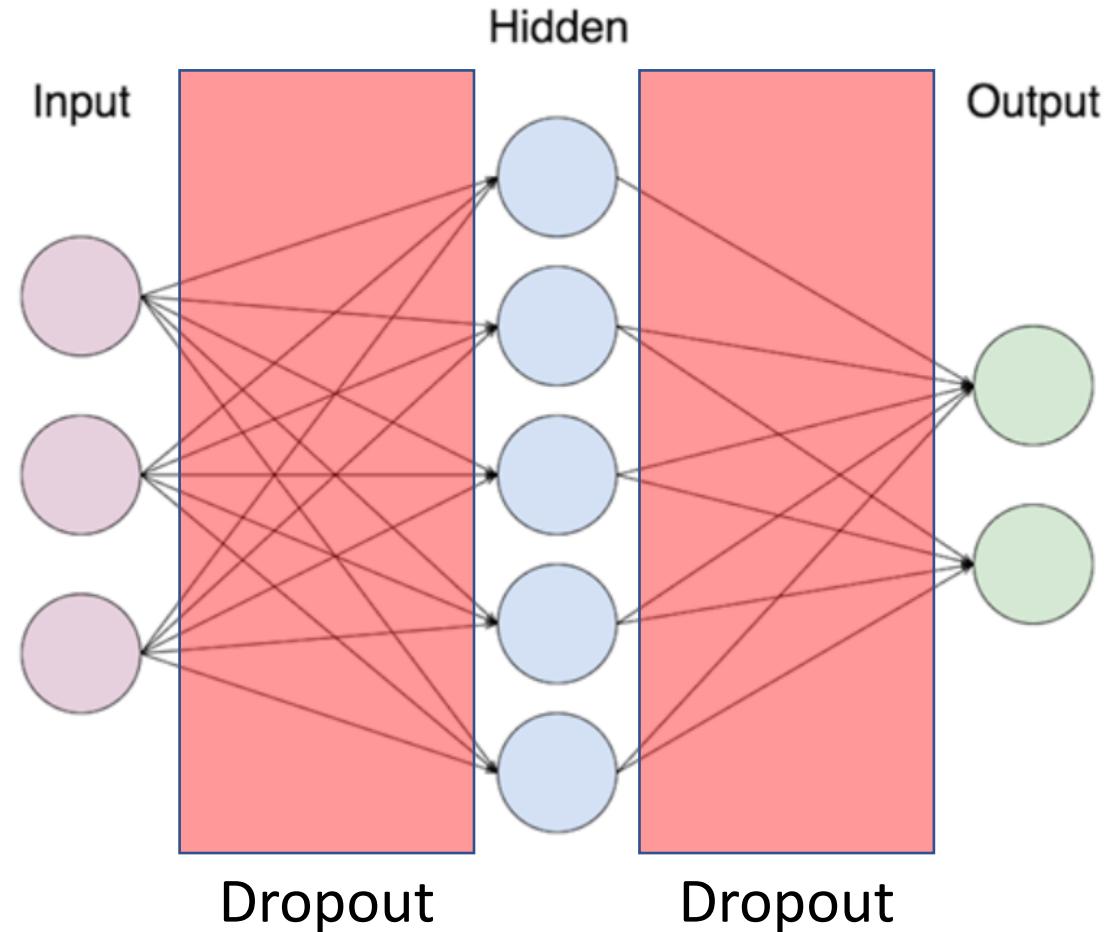
Input Normalization

Other normalization procedures include:

- Layer Norm: Transposes Batch Norm. Normalizes over all summed inputs to a layer
 - <https://arxiv.org/abs/1607.06450>
- Group Norm: Normalizes by grouped channels instead of batches
 - <https://arxiv.org/abs/1803.08494>

Pytorch Operators/Layers – Dropout Regularization

- Activation Functions
- Normalization
- Dropout

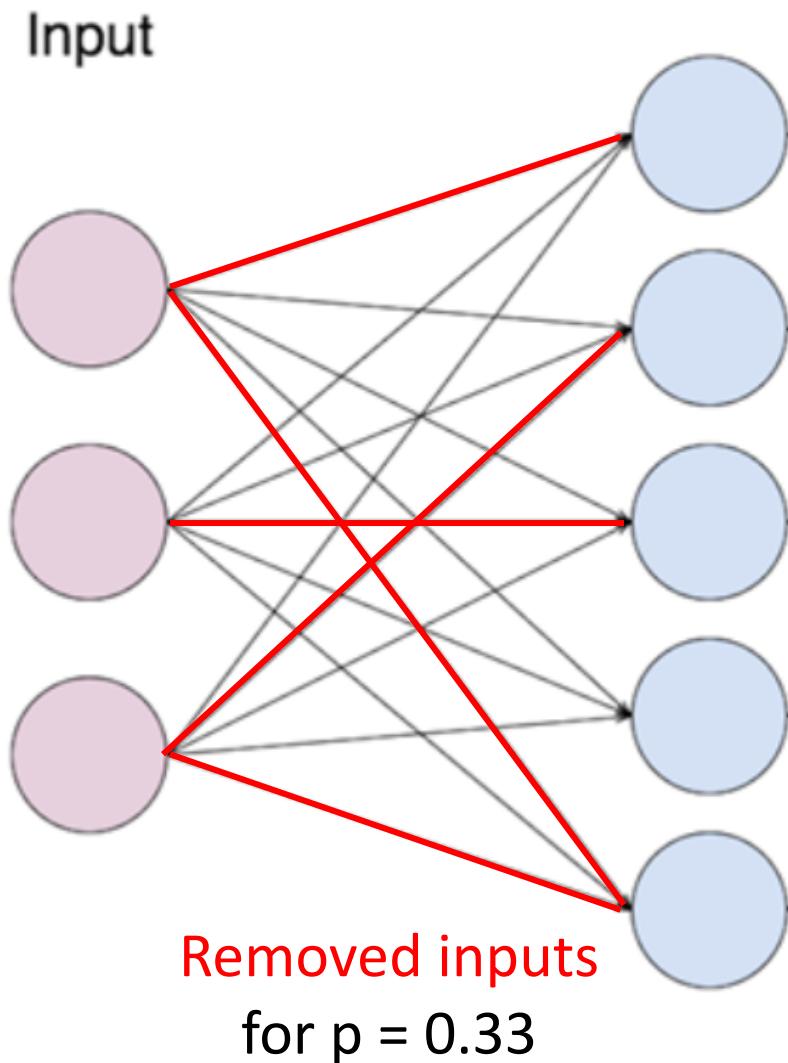


Dropout

- Randomly zeroes some elements of input tensor with probability p
- Effective technique for regularization
- Outputs scaled by $1/(1-p)$
- Treated as identity during evaluation

Syntax Example:

```
self.dp2 = torch.nn.Dropout(p=0.33)
self.layer2 = torch.nn.Linear(input_dim, output_dim)
```

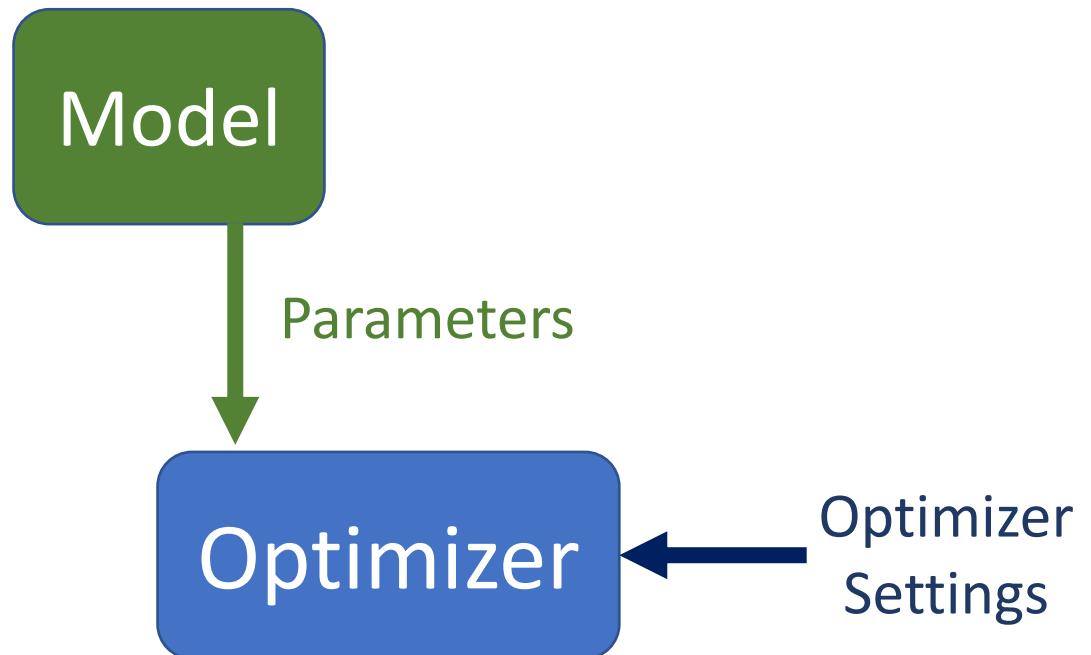


Designing Training Procedures

- Optimizer Setup
- Weight Initialization

Optimizer Initialization

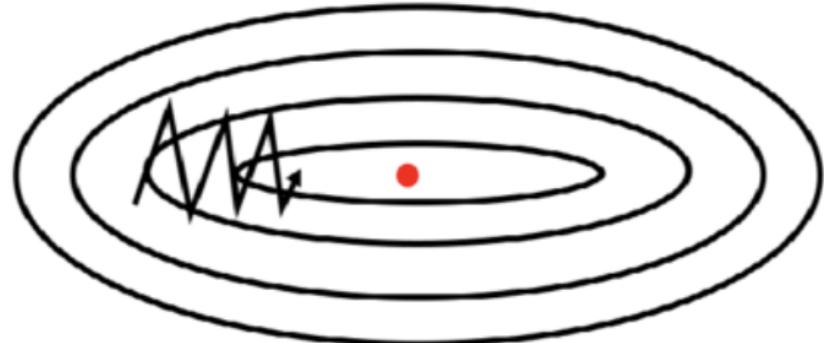
- Parameters:
 - Should be iterable containing parameters to optimize
 - E.g., `model.parameters()` or `[var1, var2]`
 - Parameters must be defined BEFORE the optimizer
- Optimizer Settings
 - Learning rate, weight decay, etc.



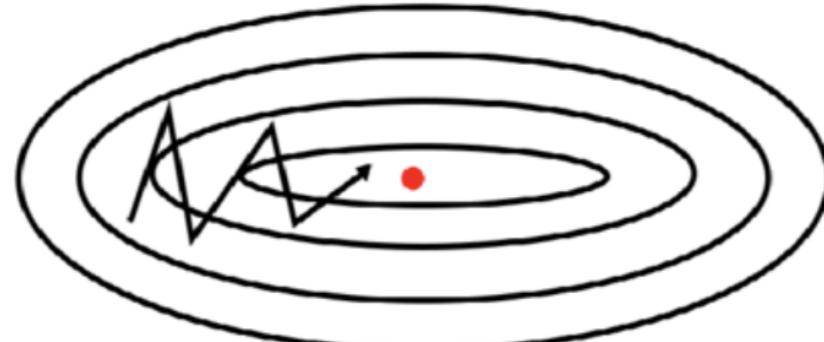
Optimizers

- Stochastic Gradient Descent
(`torch.optim.SGD`)
 - `params`: Model parameters
 - `lr`: Learning rate (required)
 - `momentum`: momentum factor (default: 0)
 - `weight_decay`: (default: 0)

SGD without momentum



SGD with momentum



Example: `torch.optim.SGD(model.parameters(), lr = 0.001, momentum = 0.2, weight_decay = 0.1)`

Optimizers

- Adam (`torch.optim.Adam`)
 - params: Model parameters
 - lr: Learning rate (default: 0.001)
 - betas: coefficients (tuple) used for computing running averages of gradient and its square (default: (0.9, 0.999))
 - eps: term added to denominator to improve numerical stability (default: 1e-8)
 - weight_decay: (default: 0)

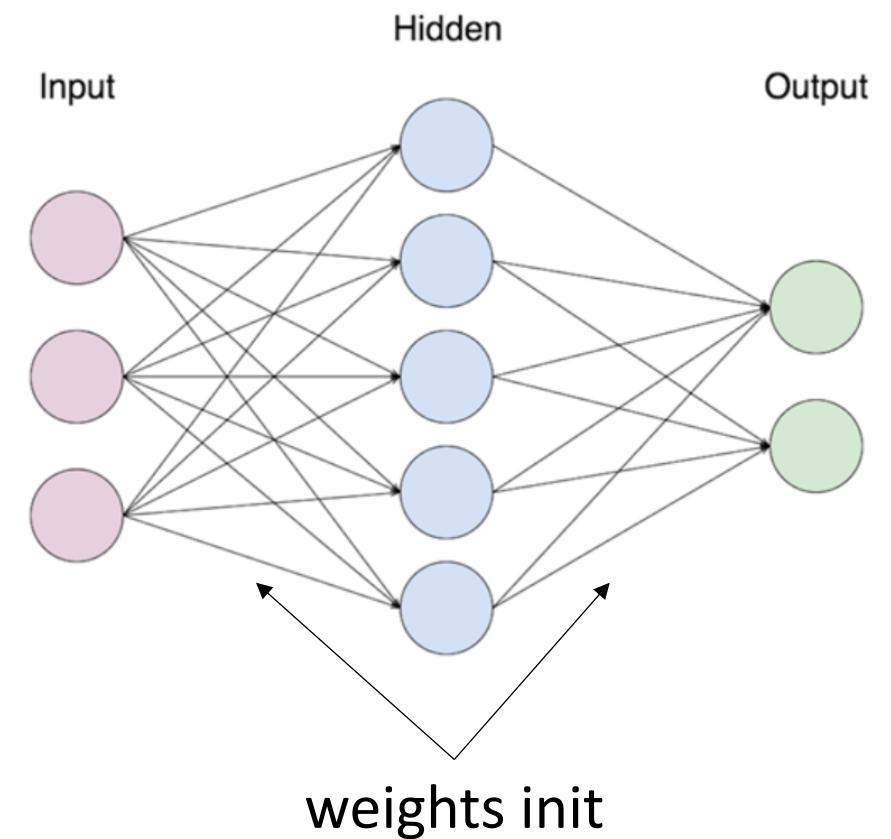
Example:

```
torch.optim.Adam(model.parameters(),  
lr = 0.01, betas = (0.95, 0.998),  
eps = 1e-7)
```

Other Common Optimizers

- AdaDelta (`torch.optim.Adadelta`)
 - Precursor to Adam which uses first-order estimates to adapt learning rate
- RMSProp (`torch.optim.RMSprop`)
 - Take the square root of the gradient average before adding epsilon to normalization of LR
- Adamax (`torch.optim.Adamax`)
 - Variant on Adam based on infinity norm

Pytorch Operators/Layers - Weight Initialization



Weight Initialization

- Initializing weights from various distributions plays a role in training. Some frequently used initialization procedures are:
 - Normal distribution
 - Xavier - compatible w. tanh
 - Kaiming (He) – compatible w. relu

Syntax Example (initialize weights of all FCN layers):

```
# define a function to init weights (here xavier_uniform)
```

```
def init_weights(m):  
    if type(m) == nn.Linear:
```

```
        torch.nn.init.xavier_uniform(m.weight)
```

```
# initialize the weights once model is created
```

```
myFCNmodel = myFCN(params)
```

```
myFCNmodel.apply(init_weights)
```

